



# Infrastructure as Code Using Terraform

...

David D. Riddle <[ddriddle@illinois.edu](mailto:ddriddle@illinois.edu)>

Jon R. Roma <[roma@illinois.edu](mailto:roma@illinois.edu)>

November 2, 2017

All original material in this presentation series is  
Copyright © 2017 by the Board of Trustees  
of the University of Illinois,  
and by the presenters

# Who are we, and why are we here?

- **David Riddle**

- David is a software developer at Technology Services

- **Jon Roma**

- Jon is a software developer who has spent 35 years at Technology Services and its predecessors

- Both David and Jon are part of Continuous Improvement team within the Application Services division of Tech Services

# Part 1

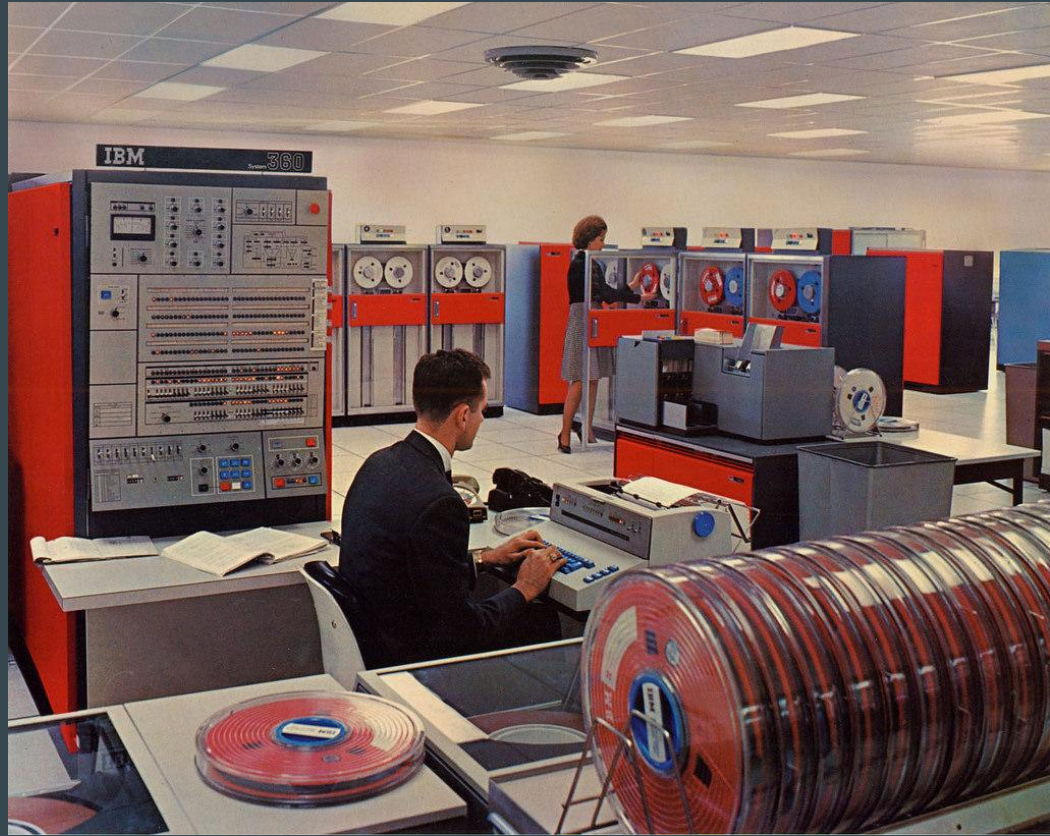
An introduction to  
Infrastructure as Code  
and to Terraform

- What is “Infrastructure as Code”?
- What is Terraform and what is the competition?
- Why did we choose Terraform over the others?

# Why infrastructure as code?

# Once upon a time...

- Infrastructure was physical
- Likewise, configuring and provisioning involved physical activity
  - Plugging in devices, and mounting media
- Nearly everything was a manual, one-off operation



# Evolving ... but not without pain

- Over time, Technology Services moved from running monolithic mainframes to “one server per service” – and then moved onward to virtualization
- Unfortunately, the resulting sprawl of servers can be difficult and costly to maintain

# Evolving ... but not without pain

- Configuring and provisioning servers was often a manual, one-off operation – just like in days of old
- Even deploying updates often involved a combination of manual checklists and *ad hoc* scripts
- In short, configuring, provisioning, testing, and deploying services could often be painful
  - As a result, we sometimes avoided making updates to software products that would have been valuable to our customers



# Many manual operations

- Resizing of services to meet demand required shuffling hardware
  - Even with virtualization, this was largely manual and hence costly in terms of human effort
  - This workload often led to poor resource utilization
    - Throwing hardware resources at the problem was an unfortunate way to save the human effort

**The missing element was  
treating infrastructure as code**

# What do we mean by infrastructure?

- What does the term “infrastructure” entail?
  - Web servers, database servers, load balancers, etc.
  - Network components (subnets, VPNs, firewall rules)
  - Persistent storage, logging, and so forth
  - ... and much more!

# What is Infrastructure as Code?

- With Infrastructure as Code (IaC), we use the same techniques that were embraced long ago for software – but now we apply them to infrastructure components
- Techniques include:
  - Version control
  - Test Driven Development (TDD)
  - Continuous Integration (CI)

# The benefits of automation

- With infrastructure automation using IaC principles, we can repeat operations across a vast swath of servers
- We never have to scratch our heads to recall the steps taken to build a particular environment
  - All the components and all the steps are documented in the infrastructure code
- The tried-and-true techniques, practices, and tools from software development are applied to infrastructure

# Improve speed, reliability, security, and quality

- Infrastructure as Code allows us to create repeatable and testable infrastructure definitions before changes are applied to business-critical systems
- We can thereby increase the speed at which changes can be made
  - We also increase the reliability, security, and quality of services
- Does infrastructure created in the AWS console improve on the old manual processes used with physical hardware?
  - **No!**

**Why did Terraform become our  
chosen Infrastructure as Code tool?**

# There are numerous tools to do Infrastructure as Code

- Ansible
- Chef
- CloudFormation
- Puppet
- Terraform



# Choosing an Infrastructure as Code tool

- Note that all the tools we just mentioned are open source, and work with multiple cloud providers – with one exception
  - CloudFormation is closed source and is AWS-only
- All are well-documented and well-supported (with enterprise support available)
- All have a substantial community of contributors and support

# Configuration management vs. orchestration

- There is some overlap in capabilities between these tools
- Ansible, Chef, and Puppet excel at configuration management
  - Their primary function is to install and manage software on existing servers
- In contrast, CloudFormation and Terraform excel at orchestration
  - Their function is to define the server infrastructure

# Docker and Packer do configuration management

- If you're using Docker and Packer, most of your configuration management needs are already satisfied
  - Docker creates containers
  - Packer creates virtual machine images
- In either case, all you need to deploy your image is a server
  - Terraform is well-suited for building immutable infrastructure

# Mutable vs. immutable infrastructure

- If you tell Ansible to install a new version of Java, the changes are made in place on your existing servers
  - This can create “configuration drift” where each server can become slightly different than the others
  - These subtle configuration differences are difficult to detect and to debug

# Mutable *vs.* immutable infrastructure

- Contrast this with an environment made up of Docker or Packer images:
  - Don't make incremental changes to existing infrastructure
    - Throw it away and redeploy!
- Any software change simply involves building a new image and deploying it on clean servers

# Mutable *vs.* immutable infrastructure

- Think of each change as a new deployment, with several benefits:
  - You know exactly what software is running on every server
    - The code you deploy is byte-for-byte identical to the code you built and tested – you've eliminated configuration drift
  - Any version of the code (past, present or future) can be deployed at will
    - This is a prerequisite for continuous integration

# Procedural vs. Declarative

- Ansible and Chef are procedural:
  - The developer writes code that defines the steps to be taken to achieve the desired end state
- Terraform, CloudFormation, and Puppet are declarative:
  - The developer writes code that defines the desired end state

# Procedural vs. Declarative

- Terraform's declarative nature means that the code always represents the desired state of your infrastructure
- It is easy at a glance to determine what's currently deployed and how it's configured
- Terraform figures out how to achieve the desired state from the infrastructure configuration code and state



# Why did Technology Services choose Terraform for IaC?

- Terraform runs on most common platforms:
  - Mac OS X
  - FreeBSD
  - Linux
  - OpenBSD
  - Solaris
  - Windows

# Why did Technology Services choose Terraform for IaC?

- Terraform supports numerous cloud providers
  - This allows using the same tool across multiple providers
- The Terraform community is enthusiastic, and the development team is prolific
  - Bugs are fixed rapidly
  - New features are added rapidly
    - Includes support for new infrastructure components

# Why did Technology Services choose Terraform for IaC?

- Infrastructure should be immutable
  - Immutable infrastructure is better suited for dynamic scaling
  - Immutable infrastructure is easier to test
- Terraform was designed from the ground up with these attributes in mind
  - Terraform rapidly gained grassroots support within several different groups in Technology Services

# Application Services' preferred technology

- **Amazon Web Services** – is our default platform for services moving to the cloud
  - Example of an exceptional case: Microsoft Azure for Windows applications like Exchange or Active Directory
- **Terraform** – Infrastructure as Code
- **Terragrunt** – a thin wrapper for Terraform supporting remote state and locking
- **Docker** – for containerizing our services

# Application Services' preferred technology

- **GitHub** – version control and issue tracking
- **Drone CI** – continuous integration engine
- **Behave, Gherkin, and Selenium** – for automated integration testing

# Application Services' implementation model

- Application Services favors pre-built Amazon machine images (AMIs) and Docker images
  - Images are faster to launch
    - Scaling up and scaling down are efficient
  - Docker images are platform-agnostic
    - They can run on any base platform (Amazon, Google Cloud, VMware, etc.)

# Application Services' implementation model

- We use Terraform to provision the infrastructure on which our containers reside
  - Our Docker clusters consist of virtual machines built from images that support Amazon EC2 Container Service (ECS)
- This means we don't need traditional configuration management tools like Ansible and Chef

# Questions?



# Part 2

In-depth view of  
Terraform concepts

- Configuration files
- Command-line options
- State
- Providers
  - Resources
  - Data sources

# Terraform configuration files

# How Terraform infrastructure definitions are organized

- Terraform manages infrastructure components by directory
  - A Terraform directory typically contains definitions for related components
  - Terraform doesn't care how the individual infrastructure definitions are organized in files within a directory
  - Furthermore, the order of definitions of resources and variables in these files doesn't matter

# How Terraform infrastructure definitions are organized

- Consider that a web server might consist of several components
- These components might include:
  - The server itself (one or more Amazon EC2 instances)
  - A load balancer
  - An autoscaler
  - Security groups (a.k.a. firewall rules)
  - A back-end database server

# How Terraform infrastructure definitions are organized

- These infrastructure components are intimately related
  - The Terraform definitions for these components – except the database tier – thus belong in a single directory
- Why might we maintain the database infrastructure separately?
  - By their very definition, databases are stateful
    - Their entire purpose is to persist valuable data
    - We ordinarily want the data to remain intact when reconfiguring the application

# How Terraform infrastructure definitions are organized

- Web servers, on the other hand, are stateless
  - We can treat these components as ephemeral infrastructure
    - We want to build, destroy, and reconfigure at will
- As a result, our general recommendation is to divide the Terraform configuration based on their statefulness
  - Give yourself some flexibility by considering your infrastructure's lifecycle

# Terraform configuration files – two file formats

- There are two accepted formats of writing Terraform infrastructure definition files:
  - HashiCorp Configuration Language (HCL)
  - JSON

# Terraform configuration files in HCL

- HCL is Terraform's native format
  - HCL is designed to strike a balance between being human-readable and machine-friendly
  - HCL is recommended format for Terraform configuration files
  - HCL configuration files are text files with names ending in **.tf**



# Terraform configuration files in JSON

- Terraform's support for JSON configuration files allows software generation of Terraform infrastructure definitions
- Terraform processes JSON files that end with `.tf.json`

# Terraform configuration files

- Any files in the working directory that do not end in either `.tf` or `.tf.json` are ignored
  - The supported formats (HCL and JSON) can be mixed freely
- We recommend sticking with HCL configuration files
  - All of our examples use HCL

# An example Terraform configuration file

```
# Define the Amazon machine image (AMI) to use.
```

```
variable "ami" {  
    description = "Amazon machine image"  
}
```

```
/* A multi  
   line comment. */
```

```
resource "aws_instance" "web" {  
    ami          = "${var.ami}"  
    count        = 2  
    instance_type = "${var.instance_type}"  
    tags {  
        Name = "example"  
    }  
}
```

- This Terraform configuration file is written in HCL
- It defines a variable named **ami**
- Second, it declares an **aws\_instance** (EC2) resource
  - The **count** parameter indicates that two instances are to be created
  - We add a tag for our instances

# Terraform command-line options

# Terraform commands (Command line interface)

## Common commands:

▶ apply	Builds or changes infrastructure
▶ console	Interactive console for Terraform interpolations
▶ destroy	Destroy Terraform-managed infrastructure
▶ env	Environment management
▶ fmt	Rewrites config files to canonical format
▶ get	Download and install modules for the configuration
▶ graph	Create a visual graph of Terraform resources
▶ import	Import existing infrastructure into Terraform
▶ init	Initialize a new or existing Terraform configuration
▶ output	Read an output from a state file
▶ plan	Generate and show an execution plan
▶ push	Upload this Terraform module to Terraform Enterprise to run
▶ refresh	Update local state file against real resources
▶ show	Inspect Terraform state or plan
▶ taint	Manually mark a resource for recreation
▶ untaint	Manually unmark a resource as tainted
▶ validate	Validates the Terraform files
▶ version	Prints the Terraform version

# The three basic Terraform commands

- The principal Terraform verbs you'll encounter are few in number
  - **plan** – Show Terraform's execution plan
    - Tells the developer what infrastructure will be created, modified, or destroyed
  - **apply** – Apply the infrastructure definition
    - Create, modify, and destroy components as required
  - **destroy** – Destroy the defined infrastructure

# Additional commonly-used Terraform commands

- The following Terraform verbs are also useful:
  - `init` – Initialize a Terraform configuration
    - Downloads and initializes any plugins (such as providers)
  - `get` – Initialize modules used in a Terraform configuration
    - Downloads module from source and initializes
  - `fmt` – Rewrites configuration files to canonical format

# Additional commonly-used Terraform commands (continued)

- Still more commonly-used Terraform verbs:
  - **show** – Inspect Terraform state or plan
  - **output** – Extract one or more outputs from a Terraform state file
    - Useful for visual confirmation of completed plan
    - Also useful to import Terraform outputs into a shell script or program



# Terraform state

# Terraform state

- Terraform maintains state about each infrastructure configuration
  - Allows tracking real resources and metadata
  - Improves performance for large infrastructure configurations
- Terraform state is maintained per directory
  - By default, the state resides in a file named `terraform.tfstate` within each directory
  - Remote Terraform state storage is better suited to team environments

# Terraform state

- Terraform uses its state to create its execution plan and to make changes to your infrastructure
  - Before any operation, Terraform synchronizes the state to reflect the real infrastructure

# Terraform state

- Terraform uses its state to map the infrastructure definitions in your Terraform files to the real world infrastructure
- You might have a resource definition like this:

```
resource "aws_instance" "web" {  
  ...  
}
```

- Terraform's state file allows mapping this resource definition to a running Amazon EC2 instance `i-006534aca5bcf2164`, for example

# Terraform state

- Terraform must also maintain metadata
  - The most obvious metadata element is dependency order
    - This is used to allow Terraform to create and destroy resources in proper sequence
  - Terraform stores additional metadata like creation time, update time, last run time, etc.
  - Cached metadata can improve performance under some conditions

# Terraform import

- We often experimentally build and configure infrastructure manually
- We eventually want Terraform to manage this as Infrastructure as Code
- Terraform can import existing infrastructure:

```
% terraform import aws_instance.web i-006534aca5bcf2164
```

# Terraform import

- Currently, Terraform's import feature has some major limitations
  - Import isn't supported for all resources
  - Terraform only imports resources into the state file
    - Import does not generate Terraform configuration code (\*.tf files) at this time

# Terraform import

- Currently, you must manually write configuration code to preserve what terraform import has stored in the state file
  - This is a little bit tedious, but `terraform plan` can help verify that the configuration code you write accurately describes the manually-created infrastructure
- A future version of Terraform will generate complete Infrastructure as Code in the form of `*.tf` files



# Terraform providers

# Terraform providers

- A key Terraform concept is that of **providers**
  - Terraform is agnostic to the underlying infrastructure platform
- Each provider is responsible for API interactions with the underlying platform
  - A provider exposes platform functionality in the form of **resources** and **data sources**
- Providers are initialized with the **terraform init** command

# Terraform has dozens of providers ... a few are listed below

- AWS
- DNS
- Docker
- GitHub
- Google Cloud
- Heroku
- HTTP
- Local
- Microsoft Azure
- MySQL
- PostgreSQL
- Random
- Template
- VMware vSphere

# Terraform providers

- This provider block defines parameters used to interact with AWS:

```
provider "aws" {  
    region = "us-east-2"  
    profile = "roma-test-readonly"  
}
```

- This provider block accomplishes two things:
  - Terraform loads the plugin supporting AWS
  - We specify the AWS region and identify the authentication credentials

# Terraform providers

- We normally don't hardcode the AWS profile in the Terraform provider block
- Instead, we set two system environment variables:
  - `AWS_DEFAULT_REGION`
  - `AWS_PROFILE`
- Then, all we need is a trivial Terraform provider block:

```
provider "aws" {}
```

# Using multiple providers is supported

```
# Provider for AWS region in North America.
provider "aws" {
  alias    = "north-america"
  region  = "ca-central-1"
}
```

```
# Provider for AWS region in Europe.
provider "aws" {
  alias    = "europe"
  region  = "eu-central-1"
}
```

NOTE: This example assumes that suitable Amazon credentials are configured

Setting the **AWS\_PROFILE** environment variable is one possible method to do this

```
# Resource for VPC in North America.
resource "aws_vpc" "vpc_na" {
  cidr_block = "10.0.0.0/16"
  provider   = "aws.north-america"
}
```

```
# Resource for VPC in Europe.
resource "aws_vpc" "vpc_europe" {
  cidr_block = "10.0.0.0/16"
  provider   = "aws.europe"
}
```

```
# Output the id from each of the VPCs.
output "vpc_na_id" {
  value = "${aws_vpc.vpc_na.id}"
}
output "vpc_europe_id" {
  value = "${aws_vpc.vpc_europe.id}"
}
```

# Terraform resources

- A **resource** usually corresponds to an infrastructure component such as:
  - Virtual machines (EC2 instances)
  - DNS records
  - Database (RDS or DynamoDB) instances
  - Security groups
  - Virtual private clouds (VPCs)
  - Lambda functions

# Some Terraform resources for AWS

- After specifying the AWS provider to Terraform, the infrastructure developer gains access to numerous AWS resources
- A few of the many available AWS resources are shown below:
  - `aws_instance` – An Amazon EC2 instance
  - `aws_s3_bucket` – An S3 bucket
  - `aws_s3_bucket_object` – An object in an S3 bucket
  - `aws_db_instance` – An RDS database instance



# Wait ... there's more!

- The above list is just a brief example
  - There are many AWS resources supported by Terraform
    - Support for new platform features in Terraform is rapid
- A resource can be almost anything – a physical machine, a VM, a network switch, a subnet, a container, etc.

# Terraform data sources

- Terraform providers also provide **data sources**, which allow retrieving data external to your Terraform configuration
  - Information not maintained in Terraform
  - Information defined by independent Terraform configurations
- Data sources present read-only views into pre-existing data, or compute values at runtime
  - Allows separate management of shared infrastructure components in isolation from consumers of those components

# Example Terraform data sources for the AWS provider

- **aws\_ami** – given a search string, get a list of Amazon machine images (AMIs) matching the search criteria
- **aws\_availability\_zones** – given an AWS region, get a list at runtime of AWS availability zones within that region
- **aws\_db\_instance** – gives information about an RDS instance
- **aws\_iam\_role** – given the name of an IAM role, get information about the role like its Amazon Resource Name (ARN)

## Example: Show available AWS availability zones in a region

```
provider "aws" {  
    region = "us-east-2"  
}  
  
data "aws_availability_zones" "available" {  
    state = "available"  
}  
  
output "az_available" {  
    value = "${data.aws_availability_zones.available.names}"  
}
```

The value in the output block is a Terraform variable reference – more about variables later

# Output from our example

```
% terraform apply
data.aws_availability_zones.available: Refreshing state...

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

Outputs:

```
az_available = [
    us-east-2a,
    us-east-2b,
    us-east-2c
]
```

# Terraform provisioners

- A **provisioner** can be used to execute scripts on a local or remote server while creating or destroying resources
  - Examples: bootstrap a resource, clean up before destroy, run configuration management, etc.
  - This provisioning might take the form of a Chef or Ansible script – or even a shell script
- Technology Services doesn't make heavy use of provisioners
  - We do provisioning at the time we build Docker images

# Terraform provisioners

- A provisioner block can be inserted inside any resource block:

```
resource "aws_instance" "web" {  
    ...  
    provisioner "local-exec" {  
        command = "echo ${self.private_ip} > ip_addr.txt"  
    }  
}
```

- Remote provisioners will need a connection block to specify connection method (e. g., ssh or winrm), username, and password

# Terraform provisioners

```
resource "aws_instance" "web" {  
  ...  
  # Copy the conf/myapp.key file to /etc/myapp.key.  
  provisioner "file" {  
    source      = "conf/myapp.key"  
    destination = "/etc/myapp.key"  
  }  
  # Copy a string's content into /tmp/myapp.log.  
  provisioner "file" {  
    content      = "ami used: ${self.ami}"  
    destination = "/tmp/myapp.log"  
  }  
  # Copy everything in config/myapp to /etc/myapp/config.  
  provisioner "file" {  
    source      = "config/myapp/"  
    destination = "/etc/myapp/config"  
  }  
}
```

- Multiple provisioners can be defined
  - They run in sequence
- Each of the three provisioners here copies data into the resource from the host running Terraform
  - The first example copies a single file
  - The second example copies text from a Terraform variable
  - The final example copies a whole directory tree



# Questions?

# Part 3

## Advanced Terraform: Resources, variables, and interpolation

- Resources and data sources
- Terraform configuration file semantics
- Declaring, assigning, and using Terraform variables
- Terraform variable interpolation

# More about Terraform resources and data sources

# Resource declarations

- In the last session, we introduced the concepts of Terraform resources and data sources
- Resources are declared as a block that identifies the infrastructure component to be managed by Terraform:

```
resource "aws_instance" "web" {  
    ami           = "${var.ami}"  
    instance_type = "${var.instance_type}"  
    count         = 2  
}
```

# Resource declarations

```
resource "aws_instance" "web" {  
    ami           = "${var.ami}"  
    instance_type = "${var.instance_type}"  
    count         = 2  
}
```

- The literal **resource** identifies the block as a resource declaration
- The quoted value **aws\_instance** identifies the resource **type**
- The quoted value **web** is used by the developer to assign a **name** to this resource instance

# Resource declarations

```
resource "aws_instance" "web" {  
  ami           = "${var.ami}"  
  instance_type = "${var.instance_type}"  
  count         = 2  
}
```

- Taken together, the **type** and **name** must be unique across your configuration
  - A Terraform configuration resides in a single directory
- The arguments used to configure a resource declaration are enclosed within the block inside braces

# Resource declarations

```
resource "aws_instance" "web" {  
  ami           = "${var.ami}"  
  instance_type = "${var.instance_type}"  
  count         = 2  
}
```

- The specific configuration arguments supported vary by resource
  - The **aws\_instance** object declared here has arguments that make sense for a virtual server, like **ami**, **instance\_type**, and **count**
  - A different resource like **aws\_s3\_bucket** would support a different set of arguments appropriate for that resource

# Resource declarations

- When **terraform apply** runs, Terraform
  - Saves the state of the new or updated infrastructure
  - Produces one or more **attributes**
    - These are akin to return values from a subroutine call
    - An example of an attribute might be a unique identifier for a virtual machine



# Resource declarations

```
resource "aws_instance" "web" {  
    ami           = "${var.ami}"  
    instance_type = "${var.instance_type}"  
    count         = 2  
}
```

- The attributes produced by Terraform are useful in interpolations
  - One infrastructure component's declaration can thus use another component's [output] attributes as input
  - This allows Terraform to identify dependencies

# Terraform configuration file semantics

# HCL configuration file semantics

- Single-line comments start with `#`
- Multi-line comments are enclosed within `/*` and `*/`
- Values are assigned as `key = value`
  - Whitespace doesn't matter
  - The RHS (value) must be an expression that evaluates to one of the valid Terraform types

# Terraform native types

- Terraform has three native types:
  - `string`
  - `list`
  - `map`

# Terraform `string` type

- String values are expressed using double quotes
  - Example: `"silly little string"`
- Multi-line strings can use shell-style “here document” syntax
  - We discourage “here documents” for the sake of clarity
    - Inline documents impede the readability of your declarations and can often be avoided

# Using “here documents” vs. reading from a file

```
resource "aws_iam_policy" "policy" {
  description = "Not recommended"
  policy = <<EOF
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "ec2:Describe*"
      ],
      "Effect": "Allow",
      "Resource": "*"
    }
  ]
}
EOF
}
```

```
resource "aws_iam_policy" "policy" {
  description = "More recommended"
  policy      = "${file("policy.json")}"
}
```

At left, we use a “here document” to assign an inline JSON policy document to our resource

At right, we use Terraform’s **file** function to specify a file containing the JSON document

# Even better is a data source to declare the policy document

```
data "aws_iam_policy_document" "policy" {
  statement {
    actions = [
      "ec2:Describe*",
    ]
    resources = ["*"]
  }
}

resource "aws_iam_policy" "policy" {
  description = "Most recommended"
  policy      = "${data.aws_iam_policy_document.policy.json}"
}
```

- The document at left uses no JSON
- Instead, a Terraform data source defines the policy document
- The policy document resource in turn is used to initialize the policy

# Terraform **string** type (continued)

- At this time, no native numeric or boolean types are available in Terraform
  - For consistent behavior, you should express numeric or boolean values as quoted strings
    - Example: **“42”** or **“false”**



# Terraform `list` type

- A list is expressed using square brackets (`[]`)
  - Example: `["foo", "bar", "baz"]`

# Terraform **map** type

- A map is expressed with braces (**{}**) and colons (**:**)
  - Example: **{ "foo": "bar", "bar": "baz" }**

# Declaring Terraform variables

# Declaring variables with Terraform

- Variables used in a Terraform declaration must be declared
- The most trivial example of a variable declaration in Terraform is

```
variable "foobar" {}
```

- The variable **foobar** is declared without specifying a type, description, or default value
  - The type thus defaults to string

# Declaring variables with Terraform

- A string variable declaration with explicit typing:

```
variable "key" {  
  type = "string"  
}
```

- A list variable declaration with implicit typing:

```
variable "aws_zones" {  
  default = ["us-east-1a", "us-east-1b"]  
}
```

# Declaring variables with Terraform

- A map variable declaration with explicit typing:

```
variable "images" {  
    type = "map"  
  
    default = {  
        us-east-1 = "image-1234"  
        us-west-2 = "image-4567"  
    }  
}
```

# Declaring variables with Terraform

- A Terraform variable declaration block can include three optional parameters:

```
variable NAME {  
    [default = default]  
    [description = description]  
    [type = type]  
}
```

# Declaring variables with Terraform

- The optional **default** parameter assigns a default value to the variable
- The optional **description** parameter provides a human-friendly description of the variable
  - This is useful when a user runs Terraform interactively
    - The description becomes part of the interactive prompt



# Declaring variables with Terraform

- The optional **type** may be any of the valid Terraform types, namely **string**, **list**, or **map**
  - If no type is specified, the type is inferred from the default
  - If no default is provided, the type is assumed to be **string**
    - If you intend a variable to be a type other than **string**, you must declare the type or set a default value
  - Terraform will fail if the variable has no value at runtime

# Terraform variable assignment

# Assigning Terraform variables

- Terraform variables can be assigned in a variety of ways:
  - From a variables file
  - From an environment variable
  - From the command line
  - From an interactive prompt
    - Only strings can be assigned interactively
- We'll give some examples in the next several slides

# Assigning Terraform variables

- If a file named **terraform.tfvars** is found in the current directory, Terraform loads it

# An example terraform.tfvars file

```
# Set values for scalars.  
is_prod = "false"  
pi      = "3.14159"
```

```
# Set values for a list.  
some_list = [  
    "1",  
    "2",  
]
```

```
# Set values for a map.  
some_map = {  
    foo = "bar"  
    bax = "qux"  
}
```

# Assigning Terraform variables

- If you want to use a different variables file, you can specify it using **-var-file** on the command line
- For example:  

```
% terraform plan -var-file foobar
```
- The **-var-file** flag can be used multiple times per command invocation

# Assigning Terraform variables

- Terraform variables can be assigned using environment variables
- The environment variables must be named `TF_VAR_name`
- For example, the following statement:

```
% TF_VAR_image=foo terraform apply
```

would set a Terraform variable named `image` to the value `foo`

- Maps and lists can be specified using environment variables in HCL format

# Assigning Terraform variables from the command line

- You can set variables from the command line with the **-var** flag:

```
% terraform plan -var access_key=foo -var secret_key=bar
```



# Assigning Terraform variables interactively

- Finally, if Terraform lacks assignment for any variables, it will prompt you interactively
  - Interactive input is only supported for string variables
  - Terraform will produce an error if any list or map variable is undefined at runtime

# Terraform output variables

# Terraform output variables

- Output variables specify the Terraform state elements that are exposed to the caller after a successful **terraform apply**
- Example output declaration:

```
output "ip" {  
    value = "${aws_eip.ip.public_ip}"  
}
```

- This defines an output variable named **ip**
  - This can be derived from interpolation, as seen above

# Terraform output variables

- Terraform displays the output when you run `terraform apply`
- Example:

```
% terraform apply
```

```
...
```

```
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

```
Outputs:
```

```
ip = 50.17.232.209
```

# Terraform output variables

- You can also query the outputs after a Terraform configuration has been applied by using the **terraform output** command:

```
% terraform output ip  
50.17.232.209
```

- This command can be used by scripts to extract outputs

# Interpolating Terraform variables

# Terraform's interpolation syntax

- Terraform's interpolation syntax is powerful
  - You can reference variables, attributes of resource definitions, outputs of module instances, and so forth
- Variable interpolations are wrapped in the construct `${}`

# Interpolation operators

- Terraform interpolations support a number of common operators
  - Math operators `+`, `-`, `*`, and `/` for floating point values
  - Math operators `+`, `-`, `*`, `/`, and `%` for integer values
  - Equality operators `==` and `!=`
  - Numerical comparison operators `>`, `<`, `>=`, and `<=`
  - Boolean operators `&&`, `||`, and unary `!`
  - The well-known ternary operator from C and other languages  
`condition ? true_value : false_value`



# Interpolating a user string variable

- Example: `${var.foo}` interpolates the value of the variable `foo`
- Next, we'll look at interpolation of each of Terraform's data types

# Interpolating a user map variable

- Example: `${var.amis["us-east-2"]}`
  - The variable `amis` is a map object
  - The variable reference interpolates the value of the `us-east-2` key within the `amis` map

# Interpolating a user list variable

- Example 1: `${var.subnets}`
  - This interpolates the subnets variable as a list
- Example 2: `${var.subnets[3]}`
  - This returns the element at the specified list index

# Interpolating attributes from a resource

- Example: `${aws_instance.web.id}`
  - The reference consists of three components:
    - The **type** of Terraform resource referenced
    - The **name** of a specific resource instance
    - The resource **attribute** to be interpolated
- Attribute values are computed by **terraform apply**

# Example: Interpolating attributes from a resource

```
resource "aws_security_group" "ssh" {  
  ingress {  
    from_port    = "22"  
    to_port      = "22"  
    protocol     = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
}
```

```
resource "aws_instance" "default" {  
  ami           = "${var.ami}"  
  instance_type = "${var.instance_type}"  
  security_groups = ["${aws_security_group.ssh.name}"]  
}
```

- We're creating an AWS security group (firewall rule) and an Amazon EC2 instance that uses that rule
- Interpolation is used to specify the Amazon machine image and instance type for the EC2 instance
- Interpolation is used to attach the security group to the EC2 instance

# Breaking down a resource interpolation reference

- Let's break down the interpolation reference we just showed:

`${aws_security_group.ssh.name}`



type



name



attribute

# Interpolating count information

- A Terraform resource block can use the `count` meta-parameter to declare multiple instances of the resource
- The “magic” interpolation `${count.index}` can be used to specify different parameters for each resource instance
  - This construct takes the place of a loop, since Terraform is declarative and not iterative

# Example: Interpolating count information

```
variable "name" {  
  default = [ "Moe", "Larry", "Curly" ]  
}
```

```
resource "aws_instance" "stooge" {  
  ami           = "${var.ami}"  
  instance_type = "${var.instance_type}"  
  count         = "${length(var.name)}"
```

```
  tags {  
    Name = "${element(var.name, count.index)}"  
    Type = "Stooge"  
  }  
}
```

In this example, we use `${count.index}` to assign each EC2 instance a unique tag



# Interpolating an attribute from a resource

- For resources declared to have multiple instances, two additional interpolations are available
  - Example 1: `${aws_instance.stooge.0.id}`
    - Returns the id attribute of the first resource
  - Example 2: `${aws_instance.stooge.*.id}`
    - Returns a list of id attributes for each resource
    - The asterisk is sometimes referred to as “splat syntax”

# Interpolating an attribute from a data source

- Example: `${data.aws_ami.debian.id}`
- The reference consists of four components:
  - The literal string `data`
  - The `type` (here, a Terraform data source)
  - The `name` of the specific data source instance
  - The `attribute` of the data source

# Example: Interpolating an attribute from a data source

```
data "aws_ami" "amazon_linux" {
  most_recent = "true"
  owners      = ["amazon"]

  filter {
    name     = "name"
    values   = ["amzn-ami-hvm-*-x86_64-gp2"]
  }
}

output "amazon_linux_id" {
  value = "${data.aws_ami.amazon_linux.id}"
}
```

- We're using a data source to select a suitable Amazon machine image
  - Runs Linux on a 64-bit x86 platform
  - Is an official Amazon image
  - Is the most recent build
- We interpolate an attribute from the data source and output it
- The interpolated value could be used in an **aws\_instance** resource

# Breaking down a resource interpolation reference

- Let's break down the interpolation reference we just showed:

`${data.aws_ami.amazon_linux.id}`

Literal  
string  
**data**

type

name

attribute

Denotes that  
this is a  
reference to a  
data source

# Interpolating conditionals

- One can conditionally configure the count:

```
resource "aws_instance" "web" {  
  ami           = "${data.aws_ami.amazon_linux.id}"  
  instance_type = "${var.instance_type}"  
  count         = "${var.define_web ? 1 : 0}"  
}
```

- The **web** resource is only configured if **`${var.define_web}`** evaluates to true
- If the VPN count is zero, the resource is not created, or is destroyed

# Interpolating conditionals

- In this example, the conditional uses the ternary operation:

```
resource "aws_instance" "web" {  
    ami          = "${data.aws_ami.amazon_linux.id}"  
    instance_type = "${var.instance_type}"  
    count        = "${var.env == "prod" ? var.prod_num : var.dev_num}"  
}
```

- The values returned by the true and false sides of the conditional must have the same data type

# Interpolation functions

- Terraform has a rich list of functions that can be used in interpolations:
  - Examples: `abs`, `basename`, `ceil`, `cidrhost`, `dirname`, `element`, `file`, `floor`, `format`, `index`, `join`, `lookup`, `max`, `min`, `pow`, `sha512`, `sort`, `split`, `substr`
- Here's a bit of sample code:

```
output "hostname" {  
  value = "${var.app}${lookup(var.suffix, var.env)}.${var.domain}"  
}
```

# Questions?



# Part 4

## Advanced Terraform

- Visualizing your infrastructure
- Terraform modules
- More Terraform interpolations
- Terraform remote state
- Adding Terragrunt to enhance our Terraform experience
- How Application Services uses Terraform and Terragrunt today
- Terraform best practices

# Visualizing your infrastructure

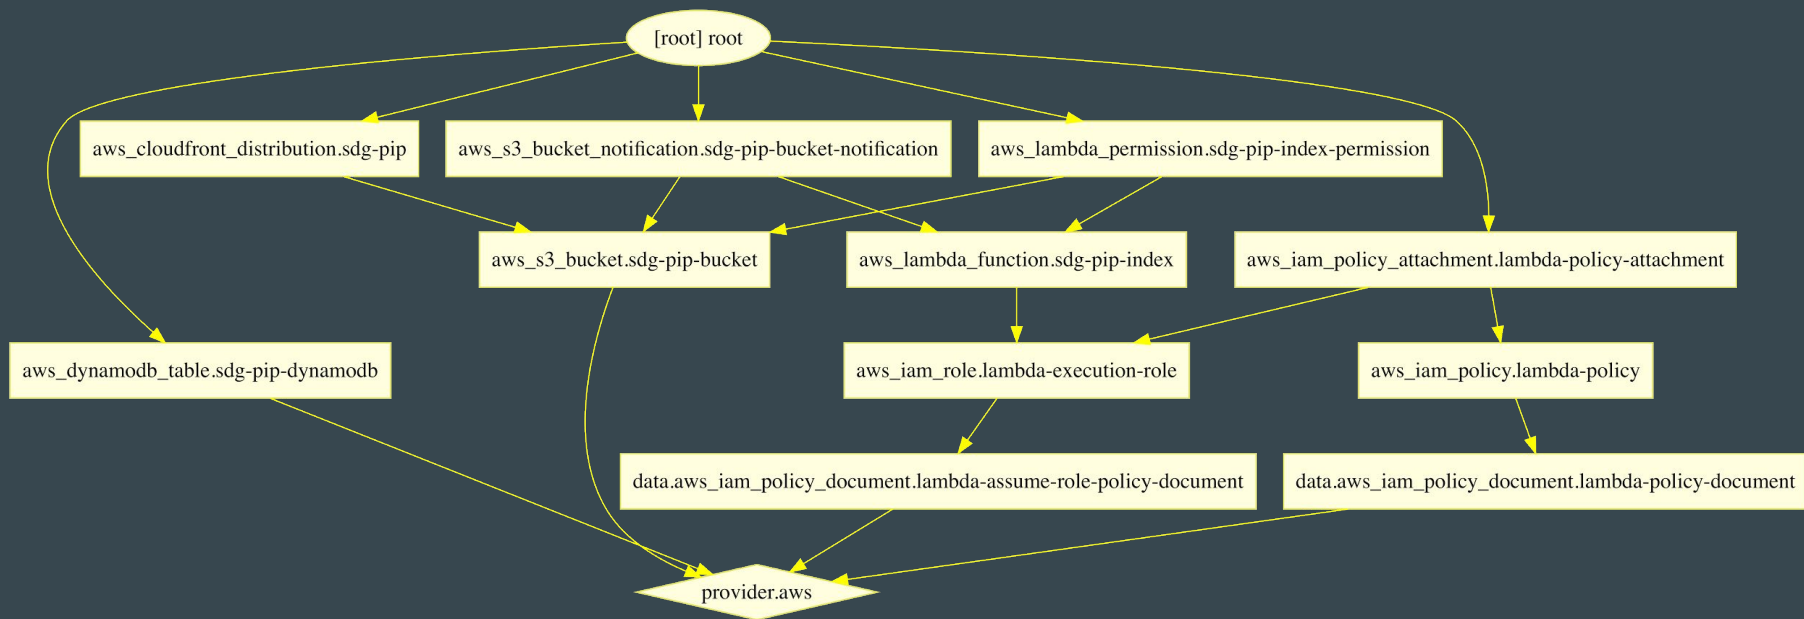
# Terraform infrastructure can be rendered into a graph

- The `terraform graph` command produces a graphical representation of a Terraform configuration or execution plan
- Usage is simple:

```
% terraform graph | dot -T png > infra.png
```

- The output is in the DOT format, and is supported by the open-source Graphviz package (see <http://www.graphviz.org/>) to generate diagrams
- A variety of web-based applications can also process DOT files

# Sample output from the **terraform graph** command



# Terraform modules

# Terraform modules – setting the stage

- Much of the Terraform documentation refers to a collection of Terraform configuration files as a module
  - Strictly speaking, a Terraform module is nothing more than a directory containing Terraform code
- Here, we'll show how we can take advantage of Terraform modules to define reusable infrastructure configuration

# Terraform modules – setting the stage

- Making a Terraform module reusable is almost effortless
  - Any directory with Terraform configuration files is inherently a Terraform module
  - Making this module reusable is simple:
    - We simply reference the module from within another Terraform configuration!
- If you refactor a Terraform module to be reusable, remove the state and variable files from the module directory – they're obsolete

# Terraform modules – setting the stage

- All our previous Terraform examples have been run from within the module's native directory
  - Unless we plan to define module variables interactively, we define their values in the **terraform.tfvars** file in the module directory
  - We've already seen that Terraform maintains the infrastructure state in the file **terraform.tfstate** in the module directory
    - This state includes any outputs defined by the developer



# Terraform modules – setting the stage

- Before we delve into invoking Terraform modules externally, we'll introduce a trivial code example
  - Our example module creates no infrastructure
  - The module defines two variables used as inputs
  - The module interpolates these two inputs, and performs exponentiation on them
  - The module produces two outputs

# A simple example of code to be used as a Terraform module

```
variable "base" {  
  description = "Base"  
}  
  
variable "exponent" {  
  description = "Exponent"  
}
```

- This example code defines variables named **base** and **exponent**
- The example code provides outputs with the names **debug\_output** and **result**

```
output "debug_output" {  
  value = "base: ${var.base} | exponent: ${var.exponent}"  
}  
  
output "result" {  
  value = "${pow(var.base,var.exponent)}"  
}
```

# Terraform modules – setting the stage

- We'll first run the module natively
  - This is the ordinary way we've used modules in the past:
    - We navigate to the module directory and run **terraform plan** and **terraform apply** as usual

# A simple example of code to be used as a Terraform module

```
% terraform apply
```

```
var.base
```

```
Base
```

```
Enter a value: 3
```

```
var.exponent
```

```
Exponent
```

```
Enter a value: 2
```

- We run the configuration *directly* here
  - Later, we'll call this code externally as a module
- We supply variable values interactively
- Terraform records state in the module directory

```
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

```
Outputs:
```

```
debug_output = base: 3 | exponent: 2
```

```
result = 9
```

# Terraform modules – invoking externally

- We've shown how the Terraform module is run natively – i. e., from its own directory
  - Now, let's invoke the module's code externally
    - In other words, let's reuse that module again and again in other Terraform configurations
- Terraform handles variable definitions and state a bit differently when the module is invoked externally

# Calling a Terraform module

- When invoking a module, Terraform manages the infrastructure defined in that module as if it were defined by the caller itself
- An important distinction is that the caller must provide any non-default variable values required by the module
- The infrastructure state (including outputs) managed by the module is not maintained in the module directory as before
  - Instead, the state produced by the module is part of the caller's state

# Calling a Terraform module

- Calling Terraform modules resembles a resource definition:

```
module "1k" {  
    base      = "10"  
    exponent  = "3"  
    source    = "github.com/cites-illinois/terraform-demo/module/power"  
}
```

- The identifier on the module declaration (“1k” in this example) denotes a name that refers to the module’s exported attributes
- The **source** identifies to Terraform where the module code resides
  - This is the only module argument required by Terraform itself

# Providing inputs to a Terraform module

- We pass two other arguments into this module:

```
module "1k" {  
  base      = "10"  
  exponent  = "3"  
  source    = "github.com/cites-illinois/terraform-demo/module/power"  
}
```

- These other arguments on the module block are simply input arguments required by the module itself
  - The module's maintainer uses **variable** declarations to determine what inputs the module accepts



# Terraform modules can reside in a variety of locations

- Terraform's modules can reside in a variety of locations:
  - GitHub repository URLs
  - Generic Git and Mercurial repository URLs
  - Amazon S3 buckets
  - Generic HTTP URLs
  - Local file paths

# Initializing module references before first use

- Before running any Terraform command using modules, you need to download the modules:

```
% terraform get
```

- This command is fast, and does not check for updates by default
  - You can thus run it repeatedly (i. e., in a CI pipeline)
  - If desired, add the **-update** flag to check for and download updates

# Receiving outputs from a module

- A module uses Terraform's **output** declarations to specify the output arguments to be returned to the caller
- We can thus interpolate a module output argument by using the format **`${module.1k.result}`**
  - The literal **module**
  - The specific module instance name **1k**
  - The attribute **result** refers to an output defined by the module

# Using multiple instances of a module

- One can use multiple instances of a module simply by providing a different name for each **module** declaration

```
module "1k" {  
    base      = "10"  
    exponent  = "3"  
    source    = "github.com/cites-illinois/terraform-demo/module/power"  
}
```

```
module "1m" {  
    base      = "10"  
    exponent  = "6"  
    source    = "github.com/cites-illinois/terraform-demo/module/power"  
}
```

# More Terraform interpolations

# Interpolating filesystem paths

- Example: `${path.module}`
  - The reference consists of two components:
    - The literal `path`
    - One of the following literal type values
      - `cwd` – This interpolates the current working directory
      - `module` – This interpolates the current module's path
      - `root` – This interpolates the root module's path

# An example of filesystem path interpolation

- Here's a trivial Terraform configuration that uses the `${path.module}` interpolation
- The code uses the current module's path to refer to a template file:

```
output "template_path" {  
    value = "${path.module}/dumb.template"  
}
```

# An example of filesystem path interpolation

- Running **terraform apply** on the above configuration might produce the following output:

```
% terraform apply
```

```
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

```
Outputs:
```

```
template_path = /var/tmp/roma/terraform-demo/part4/path/dumb.template
```



# Local variables

- Terraform's local variables are analogous to a function's local variables
  - A local variable assigns a name to an expression
    - Useful for complex expressions evaluated repeatedly
- Local variable names must be unique throughout a module
- The value can be any valid Terraform expression

# Assigning values to local variables

- Local variables are assigned values using a **locals** block:

```
locals {  
    fqdn = "${hostname}${suffix}.${domain}"  
}
```

- Each **locals** block can declare as many local variables as needed
- A module may contain any number of **locals** blocks
  - The Terraform documentation recommends grouping related local variable declarations together

# Interpolating local variables

- Interpolation of local variables takes place via a reference like `${local.fqdn}`
  - The reference consists of two components:
    - The literal `local`
    - The variable name (`fqdn` in this example)

# Interpolating local variables

```
locals {  
  gigabyte = "${floor(pow(10, 9))}"  
  terabyte = "${floor(pow(10, 12))}"  
}
```

```
output "gigabyte" {  
  value = "${local.gigabyte}"  
}
```

```
output "terabyte" {  
  value = "${local.terabyte}"  
}
```

- Within the locals block, we declare two local variables
  - Each of these variables is assigned the result of an interpolation
- We then define outputs that in turn interpolate the local variables

# Interpolating local variables

```
% terraform apply
```

```
Apply complete! Resources: 0 added, 0 changed, 0 destroyed.
```

```
Outputs:
```

```
gigabyte = 1000000000
```

```
terabyte = 1000000000000
```

# Terraform remote state

# Remote state

- As we mentioned earlier in this series, Terraform by default stores state in a file named **terraform.tfstate** in the working directory
  - This is not well-suited to multiple people doing development and testing on multiple workstations

# Remote state

- Terraform supports remote state storage as a means of maintaining consistent Terraform configurations in a multi-developer environment
  - Terraform supports several remote state backends, including AWS S3, Microsoft Azure, and Google Cloud
  - Many remote state backends support locking, which mitigates problems in a multi-developer environment



# Backends are configured within a **terraform** block

```
terraform {  
  backend "s3" {  
    bucket      = "my_bucket"  
    key         = "path/to/my/key"  
    region      = "us-east-2"  
    encrypt     = true  
    dynamodb_table = "my_lock_table"  
  }  
}
```

- This block specifies the Amazon S3 backend for remote state
- Here, we configure the location of the remote state in Amazon S3
- We want our remote state to be encrypted
- Specifying **dynamodb\_table** enables state locking with Amazon's DynamoDB

# Terraform backend configuration

- Run `terraform init` when configuring a backend for the first time
- Terraform can easily migrate local state to remote state and vice-versa

# Adding Terragrunt to enhance our Terraform experience

# Advantages of using Terragrunt to wrap Terraform

- Application Services uses Terragrunt on top of Terraform
  - <https://github.com/gruntwork-io/terragrunt>
- Terragrunt is a thin wrapper used to keep Terraform configurations DRY<sup>†</sup>
  - Terragrunt suits Application Services well because we manage dozens of Terraform modules
  - Our multi-developer workflow demands remote state

<sup>†</sup> DRY = Don't repeat yourself!

# Advantages of using Terragrunt to wrap Terraform

- Auto init (runs `terraform init` when necessary)
- Can run Terraform commands on multiple infrastructure modules:
  - Terragrunt adds new verbs `plan-all`, `apply-all`, `output-all`, and `destroy-all`
- Terragrunt allows enforcing dependencies between modules
- Terragrunt simplifies configuration of Terraform's built-in remote state and locking

# How Application Services uses Terraform and Terragrunt today

# Infrastructure hierarchy

- The Terraform infrastructure definitions for each product reside in their own GitHub repository
- We maintain a separate directory hierarchy in GitHub for deploying the various products' infrastructure components
  - Top level nodes are environments (**prod**, **qa**, **test**, etc.)
    - Within each environment are nodes with configuration for individual Terraform modules
- We therefore separate the modules from their environment-specific configuration

# Modules defining infrastructure reside in their own repositories

```
foo_app/  
└── main.tf
```

```
deploy  
├── prod/  
│   ├── foo_service/  
│   │   ├── provider.tf -> ../../provider.tf  
│   │   └── terraform.tfvars  
│   ├── provider.tf  
│   ├── terraform.tfvars  
│   └── test/  
│       └── foo_service/  
│           ├── provider.tf -> ../../provider.tf  
│           └── terraform.tfvars
```

- The node **foo\_app** represents a Git repository
  - This repository contains infrastructure definitions for **foo\_app**, which is an application



# Deployment configuration resides in a hierarchical repository

```
foo_app/  
└── main.tf
```

```
deploy  
├── prod/  
│   ├── foo_service/  
│   │   ├── provider.tf -> ../../provider.tf  
│   │   └── terraform.tfvars  
│   ├── provider.tf  
│   └── terraform.tfvars  
├── test/  
│   ├── foo_service/  
│   │   ├── provider.tf -> ../../provider.tf  
│   │   └── terraform.tfvars  
└──
```

- The **deploy** repository is a base for deploying services, and is used to maintain versioned infrastructure configuration data
- It contains subdirectories for the **prod** and **test** environments
- Below the **prod** and **test** environment subdirectories are found additional directories for configuration of individual services

# Deployment configuration resides in a hierarchical repository

```
foo_app/
└── main.tf

deploy
├── prod/
│   ├── foo_service/
│   │   ├── provider.tf -> ../../provider.tf
│   │   └── terraform.tfvars
│   ├── provider.tf
│   └── terraform.tfvars
└── test/
    ├── foo_service/
    │   ├── provider.tf -> ../../provider.tf
    │   └── terraform.tfvars
```

- The top-level **deploy** directory contains two configuration files, **provider.tf** and **terraform.tfvars**
- These files apply globally to all infrastructure managed within the **deploy** hierarchy
  - We'll come back to these files in a few minutes

# Deployment configuration resides in a hierarchical repository

```
foo_app/
└── main.tf

deploy
├── prod/
│   └── foo_service/
│       ├── provider.tf -> ../../provider.tf
│       └── terraform.tfvars
├── provider.tf
├── terraform.tfvars
├── test/
└── foo_service/
    ├── provider.tf -> ../../provider.tf
    └── terraform.tfvars
```

- Beneath the **prod** and **test** subdirectories is one directory per service – for example, **foo\_service**
- The minimal configuration code in these directories is used by Terragrunt to find the infrastructure module, and to configure it with appropriate per-environment settings

# A look at the individual elements within the deploy hierarchy

- We've had a brief look at the deployment hierarchy
  - Now, let's look at the key elements

# Top-level terraform.tfvars file

```
foo_app/
├── main.tf

deploy
├── prod/
│   ├── foo_service/
│   │   ├── provider.tf -> ../../provider.tf
│   │   └── terraform.tfvars
│   ├── provider.tf
│   └── terraform.tfvars
├── test/
│   ├── foo_service/
│   │   ├── provider.tf -> ../../provider.tf
│   │   └── terraform.tfvars
```

- Let's look first at the **terraform.tfvars** file at the top level directory of the **deploy** hierarchy

# Top-level `terraform.tfvars` file

```
terragrunt = {  
  remote_state {  
    backend = "s3"  
    config {  
      bucket      = "terraform-demo-terragrunt"  
      key         = "${path_relative_to_include()}/terraform.tfstate"  
      region      = "us-east-2"  
      encrypt     = true  
      dynamodb_table = "terraform-demo-terragrunt"  
    }  
  }  
}
```

- The `terraform.tfvars` file at the top of the `deploy` tree contains the Terragrunt configuration shared by the whole deploy tree
  - It defines the Terraform backend configuration, including the Amazon S3 location, DynamoDB table name, and encryption options

# Top-level `provider.tf` file

```
foo_app/
└── main.tf

deploy
├── prod/
│   ├── foo_service/
│   │   ├── provider.tf -> ../../provider.tf
│   │   └── terraform.tfvars
│   └── provider.tf
├── terraform.tfvars
├── test/
│   └── foo_service/
│       ├── provider.tf -> ../../provider.tf
│       └── terraform.tfvars
```

- We'll now examine the `provider.tf` file at the top level directory of the `deploy` tree
  - This file is referenced through symbolic links in the individual service directories at the bottom of the file hierarchy

# Top-level `provider.tf` file

```
provider "aws" {  
    region = "us-east-2"  
}
```

```
terraform {  
    backend "s3" {}  
}
```

- This is the `provider.tf` file in the top level of the `deploy` hierarchy
- Terragrunt fills in the backend information here from the top-level `terraform.tfvars` file
  - These two files are side-by-side in the `deploy` hierarchy



# The `terraform.tfvars` file in the bottom-level directory

```
foo_app/
└── main.tf

deploy
├── prod/
│   └── foo_service/
│       ├── provider.tf -> ../../provider.tf
│       └── terraform.tfvars
├── provider.tf
├── terraform.tfvars
└── test/
    └── foo_service/
        ├── provider.tf -> ../../provider.tf
        └── terraform.tfvars
```

- Last but not least, we examine the `terraform.tfvars` file in the `deploy/test/foo_service` directory

# The `terraform.tfvars` file in the bottom-level directory

```
terragrunt = {  
  include {  
    path = "${find_in_parent_folders()}"  
  }  
}
```

```
terraform {  
  source = ".../app/foo_app?ref=develop"  
}
```

```
# Specify module configuration.  
ami          = "ami-c5062ba0"  
instance_count = "1"  
instance_type  = "t2.nano"
```

- This `terraform.tfvars` file resides in `deploy/test/foo_service`
- It provides the configuration values that are specific to the instance of `foo_service` to be deployed into the `test` environment
- A counterpart file can be found below the `prod` environment directory as well
  - The other `terraform.tfvars` file likely has different configuration values

# The Application Services configuration is a work in progress

- Early on, Application Services implemented a configuration hierarchy for our core infrastructure that we manage with Terraform and Terragrunt
  - Terraform did not initially have support for remote state and locking
- Terraform and Terragrunt have evolved, and new capabilities have been added since our first design
  - We recognize shortcomings in our initial design
  - We see opportunities for improvement based on what we've learned
- Further research is required – stay tuned!

# Terraform best practices

# Terraform best practices (with an emphasis on AWS)

- Do use remote state with locking
  - Application Services uses
    - Amazon S3 for state
    - Amazon DynamoDB for locking
  - Do enable versioning on the Amazon S3 bucket

# Terraform best practices (with an emphasis on AWS)

- Do put your Terraform configuration files in version control
- Do use the same version control best practices for your infrastructure configurations as with your applications
  - Because it's important, we offer some suggestions about version control best practices
    - Because this subject is off-topic to this discussion, we make those suggestions in an appendix

# Terraform best practices (with an emphasis on AWS)

- Don't repeat yourself – use shared modules
  - Do reuse code in “cookie cutter” style in Terraform modules for consistency and reliability
  - Do create a **README.md** (or equivalent) for each module
    - Do describe what the module does
    - Do enumerate the meaning of input and output variables

# Terraform best practices (with an emphasis on AWS)

- Don't needlessly hardcode values into your configuration files
  - Interpolate data sources and resources where possible
    - Using data from the remote state configured by other Terraform modules is another possibility
      - This uses Terraform data sources
- Do use `terraform fmt` to make your configuration files neat



# Terraform best practices (with an emphasis on AWS)

- Do use **required\_version** to specify constraints on the Terraform version
  - Terraform is under intensive development with frequent updates
- Do likewise constrain a Terraform provider via the **version** argument in the provider block

```
terraform {  
  required_version = ">= 0.10.0"  
}
```

```
provider "aws" {  
  version = "~> 0.1"  
  region  = "${var.region}"  
}
```

The **required\_version** feature is well-hidden, but is explained at <https://www.terraform.io/docs/configuration/terraform.html>.

# Terraform best practices (with an emphasis on AWS)

- If you use multiple AWS accounts, do use the keyword `allowed_account_ids` in the provider block

```
provider "aws" {  
  region = "us-east-2"  
  allowed_account_ids = ["0123456789"]  
}
```

- Terraform protects you from building, modifying, or removing resources from the wrong Amazon account

```
% terraform plan
```

-----

```
Error running plan: 1 error(s) occurred:
```

```
* provider.aws: Account ID not allowed (9876543210)
```

# Questions?

# Demo code and other resources

We have set up a public repository on GitHub that includes:

- Demonstration code, including most of the examples given in this presentation
- Links to Terraform and Terragrunt distributions and documentation
- Links to useful resources about using these tools



# The End

## Thank you for your time!

# Appendix: Version control best practices

# Version control best practices

- Do use Git
  - In other well-known version control systems, branching and merging were feared and therefore seldom-used
- Don't forget that, with Git, a commit is local-only
  - Do remember to push your code to a remote repository

- **Do** branch and merge frequently and fearlessly
  - The ease of branching and merging is a major strength of Git
- **Do** use small, focused commits
  - Don't mix unrelated changes in the same commit
  - Each bug fix, feature addition, etc. can (and often should) be its own branch
  - Most branches should be short-lived!



# Version control best practices

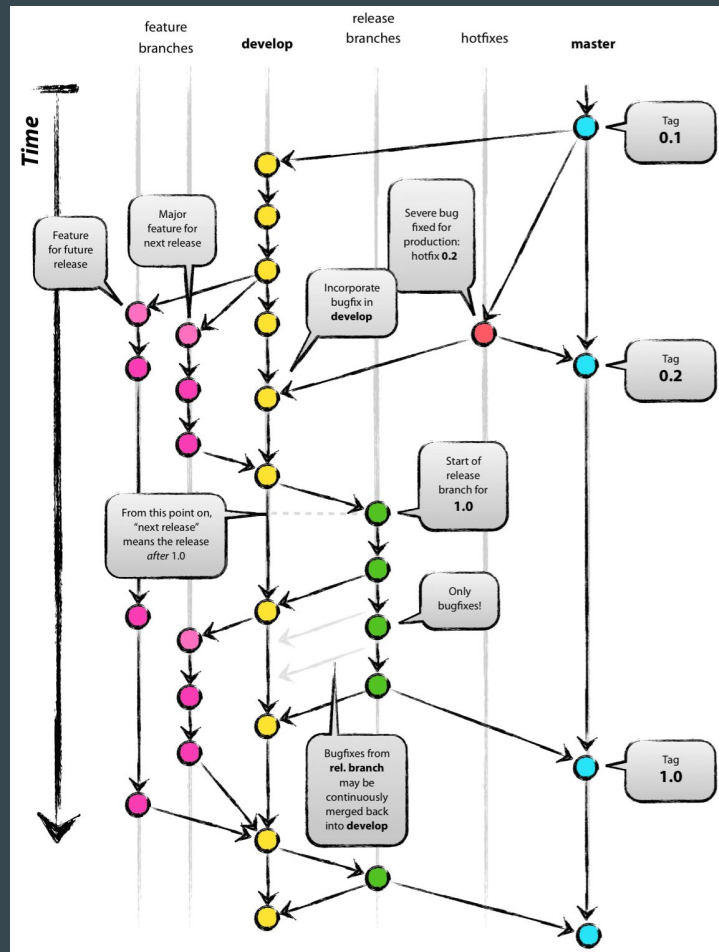
- Do commit and push often
  - Doing so helps collaboration and integration
  - Doing so helps resolve conflicts while they are still bite-sized
- Do avoid committing unfinished or non-working code
  - This can be troublesome in a continuous integration workflow
  - Do use a private branch or use **git stash** if you need to save your work in progress

# Version control best practices

- Do write meaningful commit messages
  - Sometimes intensive development results in poor commit message
  - Do get to know the **git rebase** command
- Do agree on a workflow (such as gitflow)
  - See <http://nvie.com/posts/a-successful-git-branching-model/>
  - There's a git extension that makes it easy to use gitflow:
    - <https://github.com/nvie/gitflow>

# Gitflow diagrammed

- See the description of gitflow at <https://github.com/nvie/gitflow>



# Version control best practices

- Do test your code regularly
  - Automated test suites can bring big wins
  - Automated deployment can make your life even better
- Do take care to back up your work
  - Git isn't a substitute for backups
- Do take time to isolate sensitive data from your application code and infrastructure definitions
  - This will also ease transitioning your infrastructure to the cloud